



Based on the resulting tree output "tree.cw.train", we can predict the credit rating of a hypothetical "median" customer

```
# Create a data frame for the hypothetical customer
median_cust = data.frame()
newdata = c(0,1,1,0,3,0,3,3,0,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3,3)
median_cust = rbind(median_cust, newdata)
colnames(median_cust) = names(cw.known)[-46]
```

We first create a data frame for the hypothetical customer, and then insert some data inside, then use the row-bind function `rbind` to combine the data frame and new data together, and the column name will be the same as `cw.known`.

```
> cust.pred = predict(tree.cw.train, median_cust, type="class")
> cust.pred
[1] 2
Levels: 1 2 3
```

This will give the levels result output

c)

After that, we can start to produce the confusion matrix for predicting the credit rating from the tree produce earlier

```
> confusion = with(cw.test, table(tree.pred, ~credit.rating))
> confusion
```

	credit.rating		
tree.pred	1	2	3
1	162	85	37
2	90	361	143
3	5	21	77

After we generated the confusion matrix, we can calculate the accuracy

```
> # calculate the accuracy
> # divide the sum of the diagonal values (top left down to bottom right) by the sum of all the values in the matrix
> sum(diag(confusion))/sum(confusion)
[1] 0.6116208
```

The accuracy rate is 61.2%

d)

To calculate the numerical value of the gain in entropy corresponding to the first split at the top of the tree, we first need to compute the entropy before the split by using the below formula:

$$Entropy = - \sum_{i=1}^n p_i \times \log_2 p_i$$

```

> # Compute the entropy before the split
> # get the count of all classes in credit.rating using the table() function
> beforeCountFreq = table(cw.train$credit.rating)
> #find the probability of each class
> beforeClassProb = beforeCountFreq/sum(beforeCountFreq)
> #calculate entropy (before split)
> beforeEntropy = -sum(beforeClassProb * log2(beforeClassProb))
> # Compute the entropy for 'functionary' feature value 0
> # functionary == 0
> countFreq0 = table(cw.train$credit.rating[cw.train$functionary == 0])
> classProb0 = countFreq0/sum(countFreq0)
> (functionaryEnt0 = -sum(classProb0 * log2(classProb0)))
[1] 1.366963
> # Compute the entropy for 'functionary' feature value 1
> # functionary == 1
> countFreq1 = table(cw.train$credit.rating[cw.train$functionary == 1])
> classProb1 = countFreq1/sum(countFreq1)
> (functionaryEnt1 = -sum(classProb1 * log2(classProb1)))
[1] 1.476765

```

e)

Fit a random forest model to the training set to try to improve prediction

```

> # Fit a random forest model to the training set to try to improve the prediction
> rf.cw.train = randomForest(as.factor(credit.rating)~., data = cw.train)
> rf.cw.train

Call:
randomForest(formula = as.factor(credit.rating) ~ ., data = cw.train)
Type of random forest: classification
Number of trees: 500
No. of variables tried at each split: 6

OOB estimate of error rate: 42.51%
Confusion matrix:
  1  2  3 class.error
1 61 165  0  0.7300885
2 38 440 25  0.1252485
3 11 178 63  0.7500000
> rf.pred = predict(rf.cw.train, cw.test[, -46])

```

f)

```

> # Produce confusion matrix after the tuning
> confusionRFTuned = with(cw.test, table(RFTuned.pred, credit.rating))
> # Calculate the accuracy rate after the tuning
> sum(diag(confusionRFTuned))/sum(confusionRFTuned)
[1] 0.5718654

```

With the prediction value, we can produce the confusion matrix between the test dataset and predicted value. Since the overall accuracy is 57.2% and only decreased by 4.0% compared to the accuracy without tuning, hence overfitting does not occur.

### Question 3:

By using default settings for `svm()` from the `e1071` package, we can fit a support vector machine to predict the credit ratings of customers using all of the other variables in the dataset.

a)

Predict the credit rating of a hypothetical "median" customer

```
> svmfit = svm(as.factor(credit.rating) ~ ., data = cw.train, kernel = "radial")
> print(svmfit)

Call:
svm(formula = as.factor(credit.rating) ~ ., data = cw.train, kernel = "radial")

Parameters:
  SVM-Type:  C-classification
 SVM-Kernel: radial
    cost:    1

Number of Support Vectors: 937
```

Report the decision value:

```
> predict(svmfit, median_cust, decision.values = TRUE)
1
2
attr(,"decision.values")
      2/1      2/3      1/3
1 1.021296 1.511396 -0.04938262
Levels: 1 2 3
```

b)

```
> # Generate the confusion matrix
> confusionSVM = with(cw.test, table(svm.pred, credit.rating))
> confusionSVM
      credit.rating
svm.pred  1    2    3
      1 109   56   22
      2 143  393  162
      3   5   18   73
> # Overall accuracy rate
> sum(diag(confusionSVM))/sum(confusionSVM)
[1] 0.5861366
```

c)

```
> summary(tune.svm(as.factor(credit.rating) ~ ., data = cw.train,
+                 kernel = "radial", cost = 10^c(0:2), gamma = 10^c(-4:-1)))

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:
  gamma cost
1e-04  100

- best performance: 0.393424

- Detailed performance results:
  gamma cost      error dispersion
1 1e-04     1 0.4872397 0.04708242
2 1e-03     1 0.4780664 0.04701287
3 1e-02     1 0.3985055 0.06441805
4 1e-01     1 0.4872397 0.04708242
5 1e-04    10 0.4648114 0.05250497
6 1e-03    10 0.3964646 0.06941579
7 1e-02    10 0.4637291 0.05645593
8 1e-01    10 0.4872397 0.04708242
9 1e-04   100 0.3934240 0.06817821
10 1e-03   100 0.3974954 0.05321052
11 1e-02   100 0.4810864 0.05626634
12 1e-01   100 0.4872397 0.04708242

> # Fit a model using svm
> svmTuned = svm(as.factor(credit.rating) ~ ., data = cw.train, kernel = "radial",
+               cost=100,
+               gamma = 0.0001)
> # Predict the values on test set
> svmTuned.pred = predict(svmTuned, cw.test[, -46])
> # Produce confusion matrix
> confusionTunedSVM = with(cw.test, table(svmTuned.pred, credit.rating))
> # Overall accuracy rate
> sum(diag(confusionTunedSVM))/sum(confusionTunedSVM)
[1] 0.6034659
```

#### Question 4:

Fit the Naive Bayes model to predict the credit ratings of customers using all of the other variables in the dataset.

```
> nb = naiveBayes(as.factor(credit.rating) ~ ., data=cw.train)
```

a)

```
> predict(nb, median_cust, type='class')
[1] 1
Levels: 1 2 3
> predict(nb, median_cust, type='raw')
      1      2      3
[1,] 0.9850729 0.01393277 0.0009942948
```

b)

We can use the Naïve Bayes model to predict, produce the confusion matrix, and then compute the accuracy rate.

```
> #predict the values on test set
> nb.pred = predict(nb, cw.test[,-46])
> #produce confusion matrix
> confusionNB = with(cw.test, table(nb.pred, credit.rating))
> confusionNB
      credit.rating
nb.pred  1    2    3
  1  252  439  173
  2    0    4    6
  3    5   24   78
> #calculate the accuracy rate
> sum(diag(confusionNB))/sum(confusionNB)
[1] 0.3404689
```

#### Question 5:

a) Which of the classifiers look to be the best?

As we can see from the above results, decision tree classifier gives an overall accuracy of 61.2% and 57.2% after tuning. SVM classifier gives an overall accuracy of 58.6% and 60.3% after tuning. Naïve Bayes classifier gives an overall accuracy of 34.0% which is too low compared to decision tree and SVM classifier. Since both decision tree and SVM classifier have only a small difference between the accuracy before and after tuning, the process maximized the model's performance without overfitting and therefore SVM classifier looks the best.

b) Are there any categories that all classifiers seem to have trouble with?

Since I have calculated the entropy before and after split, I realized that the entropy gained from 1.37 to 1.48 after the entropy split, which means by splitting the "functionary" column, our entropy increases, hence the functionary category has a high level of disorder and it is not suitable to use it as a category for any classifier for training and testing the data set.

## Question 6:

a)

```
# Create a logit model using all of the other variables in the dataset, with no interactions
> glm.fit <- glm(as.factor((credit.rating==1))~., data = cw.train, family = binomial)
> glm.fit

Call:  glm(formula = as.factor((credit.rating == 1)) ~ ., family = binomial,
          data = cw.train)

Coefficients:
              (Intercept)
             -17.551605
re.balanced..paid.back..a.recently.overdrawn.current.account
              1.501222
gender
             0.577104
credit.refused.in.past.
             -0.935877
savings.on.other.accounts
             -0.548195
max..account.balance.12.months.ago
             -0.004444
avrg..account.balance.12.months.ago
             0.124652
min..account.balance.11.months.ago
             -0.110469
max..account.balance.10.months.ago
             0.019305
avrg..account.balance.10.months.ago
             -0.050933
min..account.balance.9.months.ago
             -0.038009
max..account.balance.8.months.ago
             -0.019017
avrg..account.balance.8.months.ago
             -0.106852
min..account.balance.7.months.ago
             -0.094176
max..account.balance.6.months.ago
             0.069171
avrg..account.balance.6.months.ago
             -0.025278
min..account.balance.5.months.ago
             -0.088221
max..account.balance.4.months.ago
             0.034718
avrg..account.balance.4.months.ago
             0.020068
min..account.balance.3.months.ago
             0.014149
max..account.balance.2.months.ago
             0.100711
avrg..account.balance.2.months.ago
             -0.038225
min..account.balance.1.months.ago
             -0.000658
functionary
             1.740533
FI30.credit.score
             16.502759
x0..accounts.at.other.banks
             -0.027413
years.employed
             0.672572
self.employed.
             -0.376394
min..account.balance.12.months.ago
             0.030192
max..account.balance.11.months.ago
             -0.010150
avrg..account.balance.11.months.ago
             0.052783
min..account.balance.10.months.ago
             -0.101696
max..account.balance.9.months.ago
             0.096730
avrg..account.balance.9.months.ago
             -0.032928
min..account.balance.8.months.ago
             -0.041455
max..account.balance.7.months.ago
             -0.018414
avrg..account.balance.7.months.ago
             -0.074021
min..account.balance.6.months.ago
             -0.033830
max..account.balance.5.months.ago
             0.015218
avrg..account.balance.5.months.ago
             -0.072089
min..account.balance.4.months.ago
             -0.036728
max..account.balance.3.months.ago
             -0.144583
avrg..account.balance.3.months.ago
             -0.010770
min..account.balance.2.months.ago
             -0.065585
max..account.balance.1.months.ago
             -0.073012
avrg..account.balance.1.months.ago
             -0.068570

Degrees of Freedom: 980 Total (i.e. Null);  935 Residual
Null Deviance:      1059
Residual Deviance: 820.8      AIC: 912.8
```

b)

Report the summary of the model

```
> options(width = 130)
> summary(glm.fit)

Call:
glm(formula = as.factor((credit.rating == 1)) ~ ., family = binomial,
    data = cw.train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.00215  -0.65353  -0.42668  -0.00012   2.70789

Coefficients:
(Intercept)
functionary
re.balanced..paid.back..a.recently.overdrawn.current.account
FI30.credit.score
gender
X0..accounts.at.other.banks
credit.refused.in.past.
years.employed
savings.on.other.accounts
self.employed.
max..account.balance.12.months.ago
min..account.balance.12.months.ago
avrg..account.balance.12.months.ago
max..account.balance.11.months.ago
min..account.balance.11.months.ago
avrg..account.balance.11.months.ago
max..account.balance.10.months.ago
min..account.balance.10.months.ago
avrg..account.balance.10.months.ago
max..account.balance.9.months.ago
min..account.balance.9.months.ago
avrg..account.balance.9.months.ago
max..account.balance.8.months.ago
min..account.balance.8.months.ago
avrg..account.balance.8.months.ago
max..account.balance.7.months.ago
min..account.balance.7.months.ago
avrg..account.balance.7.months.ago
max..account.balance.6.months.ago
min..account.balance.6.months.ago
avrg..account.balance.6.months.ago
max..account.balance.5.months.ago
min..account.balance.5.months.ago
avrg..account.balance.5.months.ago
max..account.balance.4.months.ago
min..account.balance.4.months.ago
avrg..account.balance.4.months.ago
max..account.balance.3.months.ago
min..account.balance.3.months.ago
avrg..account.balance.3.months.ago
max..account.balance.2.months.ago
min..account.balance.2.months.ago
avrg..account.balance.2.months.ago
max..account.balance.1.months.ago
min..account.balance.1.months.ago
avrg..account.balance.1.months.ago
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1058.95  on 980  degrees of freedom
Residual deviance:  820.79  on 935  degrees of freedom
AIC: 912.79

Number of Fisher Scoring iterations: 16

      Estimate Std. Error z value Pr(>|z|)
(Intercept) -17.551605 429.995589 -0.041 0.96744
functionary  1.740533  0.183036  9.509 < 2e-16 ***
re.balanced..paid.back..a.recently.overdrawn.current.account 1.501222  0.550965  2.725 0.00644 **
FI30.credit.score 16.502759 429.993845  0.038 0.96939
gender  0.577104  0.178807  3.228 0.00125 **
X0..accounts.at.other.banks -0.027413  0.063141 -0.434 0.66417
credit.refused.in.past. -0.935877  0.341848 -2.738 0.00619 **
years.employed  0.672572  0.269126  2.499 0.01245 *
savings.on.other.accounts -0.548195  0.204670 -2.678 0.00740 **
self.employed. -0.376394  0.236506 -1.591 0.11150
max..account.balance.12.months.ago -0.004444  0.062647 -0.071 0.94345
min..account.balance.12.months.ago  0.030192  0.063737  0.474 0.63572
avrg..account.balance.12.months.ago  0.124651  0.065028  1.917 0.05525 .
max..account.balance.11.months.ago -0.010150  0.063924 -0.159 0.87385
min..account.balance.11.months.ago -0.110469  0.064328 -1.717 0.08593 .
avrg..account.balance.11.months.ago  0.052783  0.065196  0.810 0.41816
max..account.balance.10.months.ago  0.019305  0.062526  0.309 0.75750
min..account.balance.10.months.ago -0.101696  0.063199 -1.609 0.10759
avrg..account.balance.10.months.ago -0.050933  0.065720 -0.775 0.43834
max..account.balance.9.months.ago  0.096730  0.062586  1.546 0.12221
min..account.balance.9.months.ago -0.038009  0.064765 -0.587 0.55728
avrg..account.balance.9.months.ago -0.032928  0.062640 -0.526 0.59912
max..account.balance.8.months.ago -0.019017  0.063459 -0.300 0.76443
min..account.balance.8.months.ago -0.041455  0.062710 -0.661 0.50858
avrg..account.balance.8.months.ago -0.106852  0.063685 -1.678 0.09338 .
max..account.balance.7.months.ago -0.018414  0.063321 -0.291 0.77120
min..account.balance.7.months.ago -0.094176  0.063702 -1.478 0.13930
avrg..account.balance.7.months.ago -0.074021  0.061950 -1.195 0.23215
max..account.balance.6.months.ago  0.069171  0.064686  1.069 0.28492
min..account.balance.6.months.ago -0.033830  0.062428 -0.542 0.58788
avrg..account.balance.6.months.ago -0.025278  0.062786 -0.403 0.68724
max..account.balance.5.months.ago  0.015218  0.061902  0.246 0.80581
min..account.balance.5.months.ago -0.088221  0.064391 -1.370 0.17066
avrg..account.balance.5.months.ago -0.072089  0.063401 -1.137 0.25553
max..account.balance.4.months.ago  0.034718  0.062889  0.552 0.58091
min..account.balance.4.months.ago -0.036728  0.064179 -0.572 0.56714
avrg..account.balance.4.months.ago  0.020068  0.063954  0.314 0.75368
max..account.balance.3.months.ago -0.144584  0.062966 -2.296 0.02166 *
min..account.balance.3.months.ago  0.014149  0.064191  0.220 0.82554
avrg..account.balance.3.months.ago -0.010770  0.064635 -0.167 0.86767
max..account.balance.2.months.ago  0.100711  0.063196  1.594 0.11102
min..account.balance.2.months.ago -0.065585  0.063059 -1.040 0.29832
avrg..account.balance.2.months.ago -0.038225  0.064392 -0.594 0.55276
max..account.balance.1.months.ago -0.073012  0.065482 -1.115 0.26486
min..account.balance.1.months.ago -0.000658  0.062229 -0.011 0.99156
avrg..account.balance.1.months.ago -0.068570  0.064302 -1.066 0.28626
```



c)

Based on the summary above, we can see “functionary” and “re.balanced..paid.back..a.recently.overdrawn.current.acount” are appear to be significant since they have the highest positive estimate rate.

However, “FI3O.credit.score” is likely to be spuriously since the estimate rate and std.error is times higher than other predictors.

d)

```
# Use SVM model to predict the training set
# Fit an SVM model of your choice to the training set
summary(tune.svm(as.factor((credit.rating==1)) ~ ., data = cw.train,
                 kernel = "radial", cost = 10^c(-2:2), gamma = 10^c(-4:1),
                 type="c"))
(svm2 = svm(I(credit.rating == 1) ~ ., data = cw.train, type = "c"))

# Predict the values on test set[SVM]
svm.fit.pred = predict(svm2, cw.test[, -46], decision.values = TRUE)

# Predict the values on test set[GLM]
glm.fit.pred = predict(glm.fit, cw.test[, -46])
```

e)

```
library(ROCR)
# Make prediction using SVM
confusionSVM = prediction(-attr(svm.fit.pred, "decision.values"),
                           cw.test$credit.rating == 1)

# Create rocs curve based on prediction
rocSVM <- performance(confusionSVM, "tpr", "fpr")

#make prediction using Logidtic Regression
confusionGLM = prediction(glm.fit.pred, cw.test$credit.rating == 1)

#create rocs curve based on prediction
rocGLM <- performance(confusionGLM, "tpr", "fpr")

# Produce ROC chart
# Plot the graph
plot(rocGLM, col=1)
plot(rocSVM, col= 2 ,add=TRUE)
abline(0, 1, lty = 3)
# Add the legend to the graph
legend(0.6, 0.6, c('glm','svm'), 1:2)
```

